



UNIVERSITY  
of VIRGINIA

# TextAttack:

## Generalizing Adversarial Examples to Natural Language Processing

@ UVA Human and Machine Intelligence Seminar  
2021/04/13

Yanjun Qi

<http://www.cs.virginia.edu/yanjun/>



UNIVERSITY  
of VIRGINIA

# TextAttack:

## Generalizing Adversarial Examples to Natural Language Processing

@ UVA Human and Machine Intelligence Seminar  
2021/04/13

Yanjun Qi

<http://www.cs.virginia.edu/yanjun/>

Background:

Natural Language Processing and Recent  
advances by Deep Learning

# What is Natural language processing (NLP)

**Wiki:** is a field of computer science, artificial intelligence, and computational linguistics concerned with the interactions between computers and human (**natural**) languages.



- Identify the **structure** and **meaning** of **words**, **sentences**, **texts** and **conversations**
- **Deep** understanding of **broad** language
- NLP is all around us

# Machine translation

A screenshot of a Google search results page. The search query "buenas noches" is entered in the search bar. Below the search bar, the "All" tab is selected, along with other options like Images, Shopping, Apps, Videos, More, and Search tools. A message indicates "About 20,800,000 results (0.54 seconds)". The main result is a machine translation card. It shows "Spanish" on the left and "English" on the right. The Spanish phrase "buenas noches" is on the left, and its English translation "Goodnight" is on the right. There are microphone and speaker icons above the text. Below the main translation, there is a button labeled "Edit" next to the Spanish text. At the bottom of the card, there is a downward arrow icon followed by the text "3 more translations". A link "Open in Google Translate" is located at the bottom right of the card.

# Dialog Systems

**Gift shop**

Items such as caps, t-shirts, sweatshirts and other miscellanea such as buttons and mouse pads have been designed. In addition, merchandise for almost all of the projects is available.



**CD or DVD**

There is a series of CDs/DVDs with selected Wikipedia content being produced by Wikipedians and SOS Children.

**Downloading**

Downloading content from Wikipedia is free of charge.

All text content is licensed under the [GNU Free Documentation License](#)

(GFDL). Images and other files are available under different terms, as detailed on

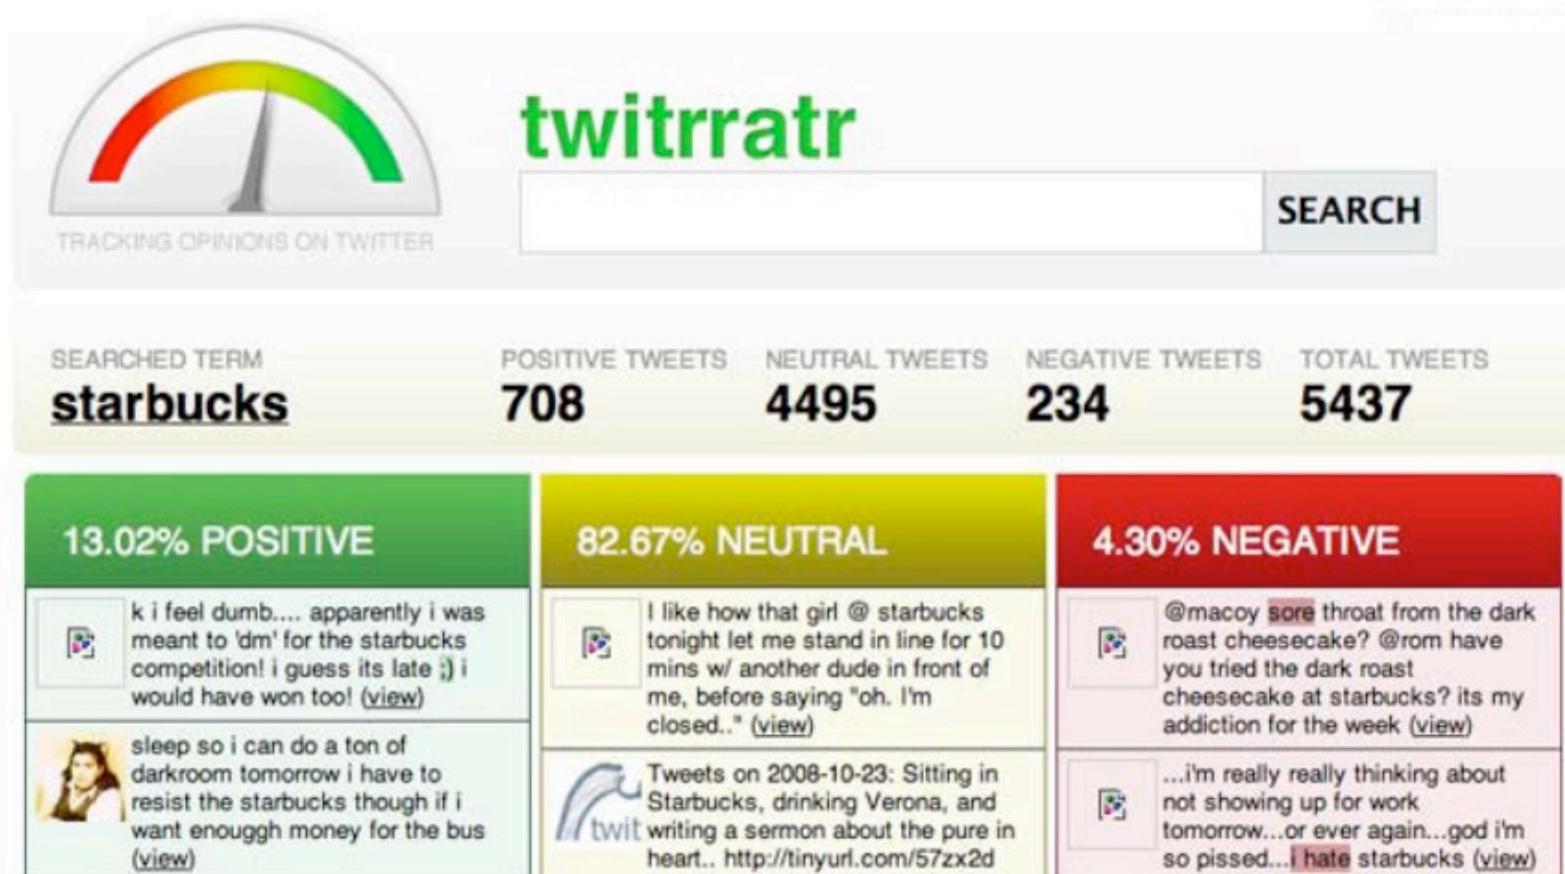
**Hi. I'm your automated online assistant. How may I help you?**

Ask

## Natural language instruction



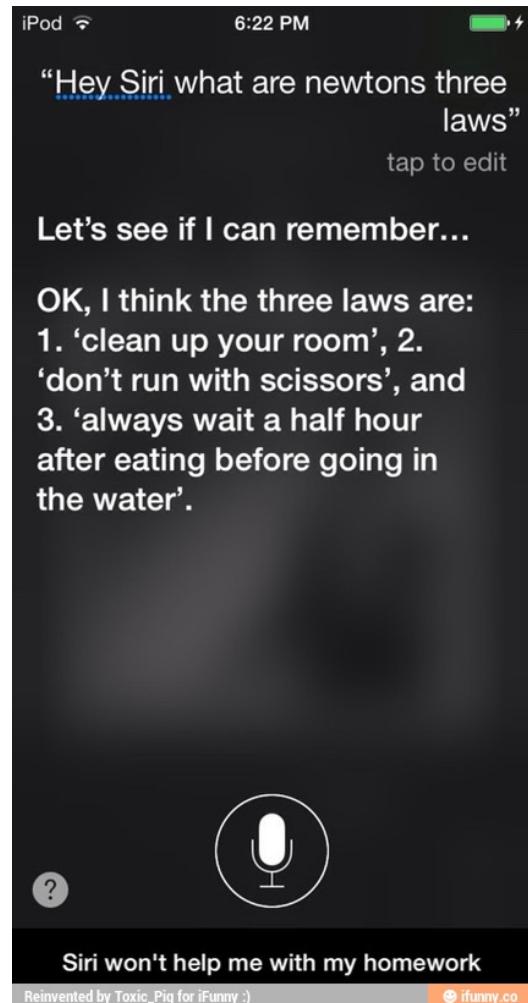
# Sentiment/Opinion Analysis



# Question answering



['Watson' computer wins at 'Jeopardy'](#)



credit: ifunny.com

# Text Classification



BIDNESS ETC

Did you mean to attach files?  
You wrote "is attached" in your message, but there are no files attached. Send anyway?

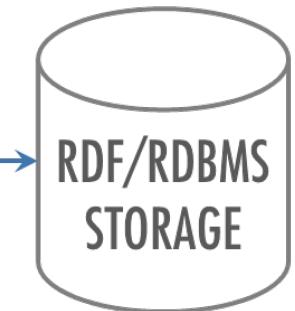
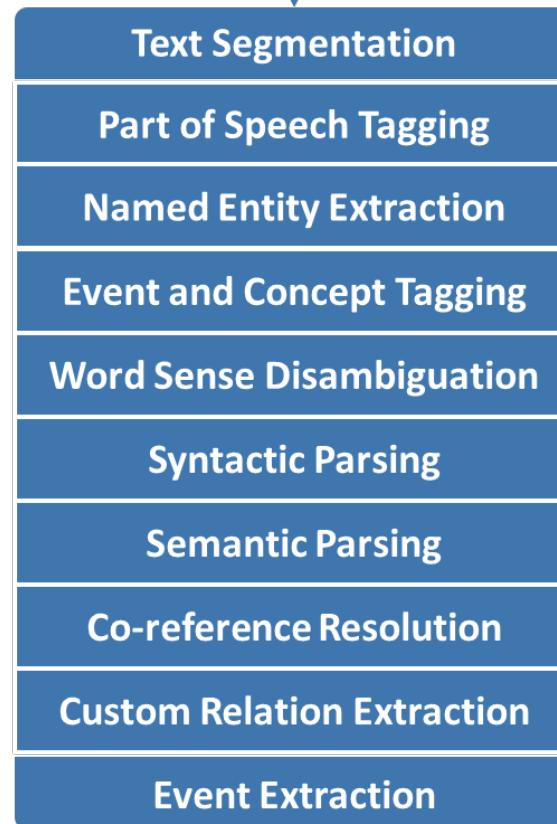
OK Cancel

1–21 of 21 < > Settings

Primary	Social Google+ 1 new	Promotions Google Offers, Zagat 2 new	Updates Google Play 1 new
James, me (2)	Hiking Hiking trip on Saturday - Yay - so glad you can join. We should leave from I	3:14 pm	
Hannah Cho	Thank you - Keri - so good that you and Steve were able to come over. Thank you :	3:05 pm	
Ian Birdsong	School Upcoming school conference dates Hello everyone. A few people have	www.wired.com	



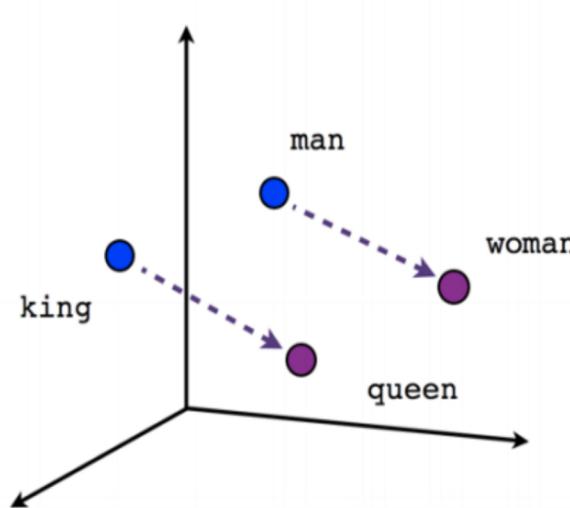
Classic NLP Pipeline  
Includes a set of  
Components for  
Understanding Text



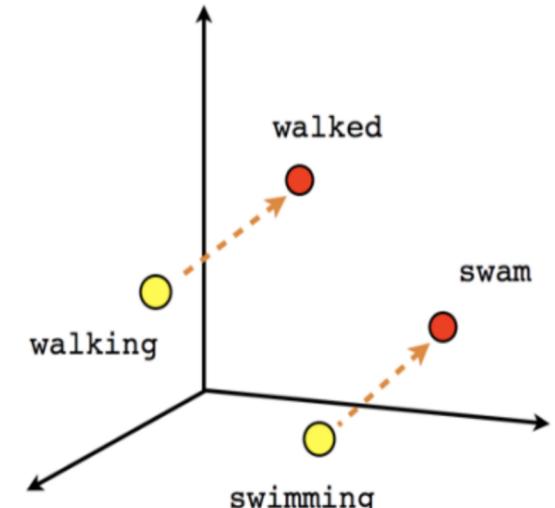
# Recent deep learning advances **on** natural language

- Before Deep NLP (Pre 2012)
  - Supervised predictors for each component
  - (BOW / LSI / Topic LDA )
- Word2Vec (2013-2016)
  - (GloVe/ FastText)
- Recurrent NN (2014-2016)
  - LSTM
  - Seq2Seq
- Attention / Self-Attention (2016 – now )
  - Attention
  - Transformer (self-attention, attention only)
  - BERT / XLNet/ GPT-2 / T5 ...

# Distributional Word Embedding Vector: To Represent A Word in DNN

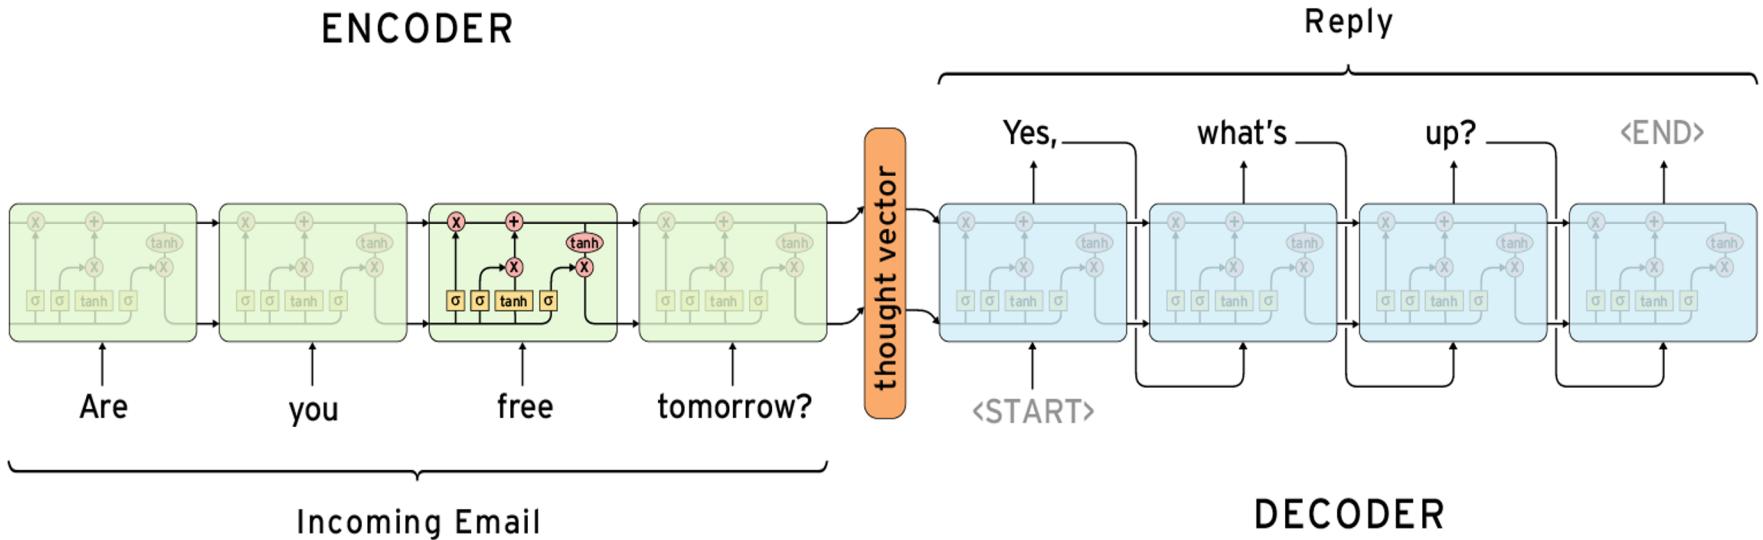
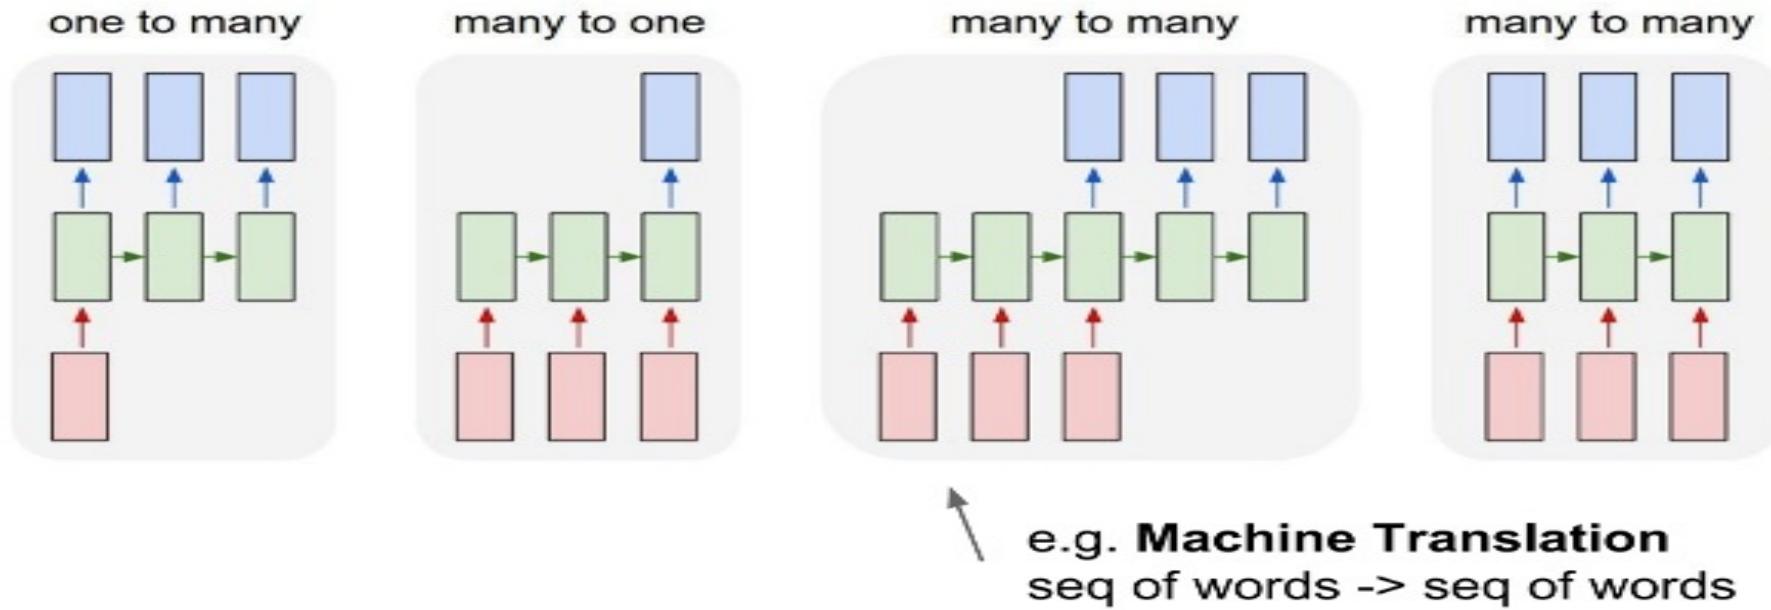


Male-Female

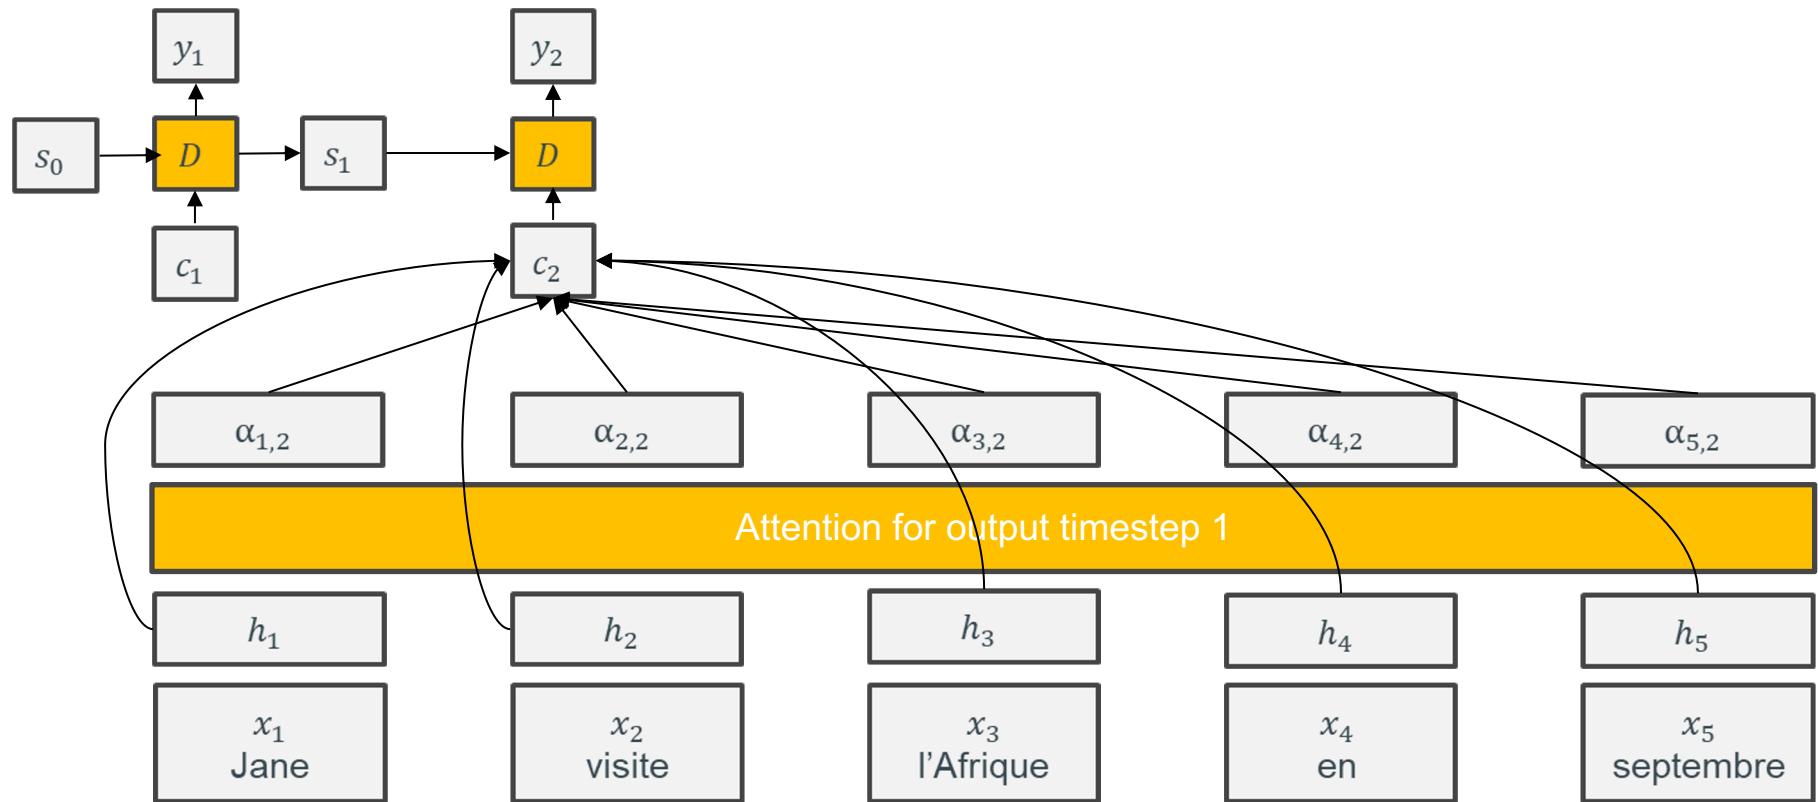


Verb tense

# Recurrent Neural Networks (RNNs) can handle



The attention module gives us a weight for each input.



Self-attention creates attention layers mapping from a sequence to itself.

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

# Transformer: Exploiting Self Attentions

- A Google Brain model.
  - Variable-length input
  - Fixed-length output (but typically extended to a variable-length output)
  - **No recurrence**
  - Surprisingly not patented.
- Uses 3 kinds of attention
  - Encoder self-attention.
  - Decoder self-attention.
  - Encoder-decoder multi-head attention.

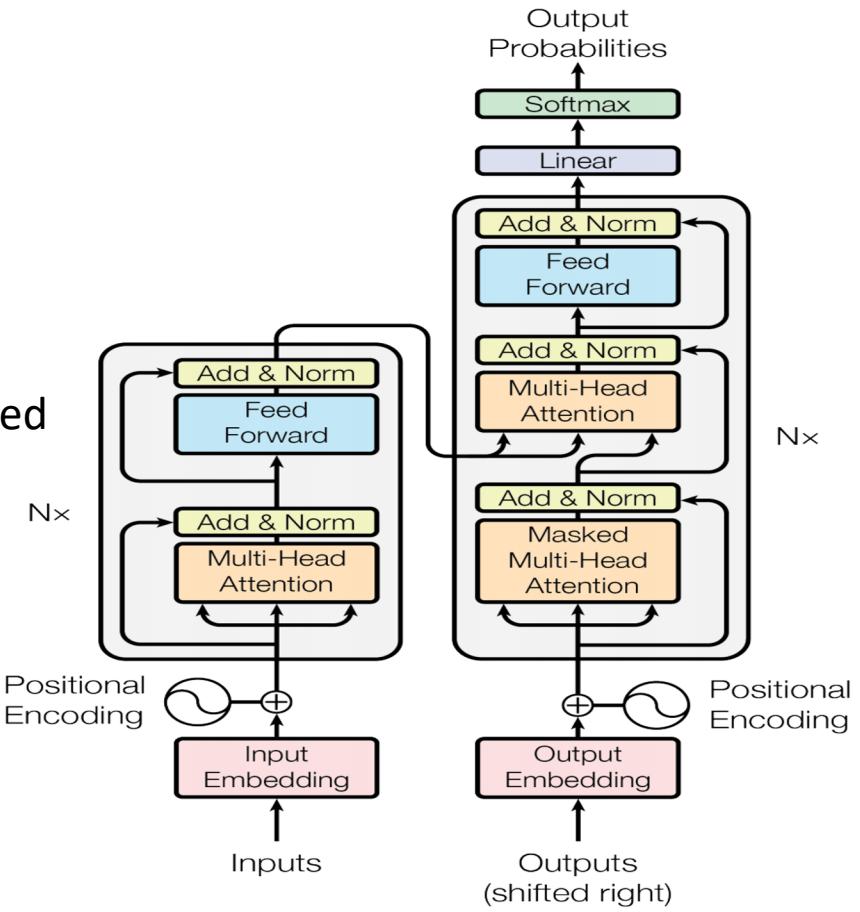


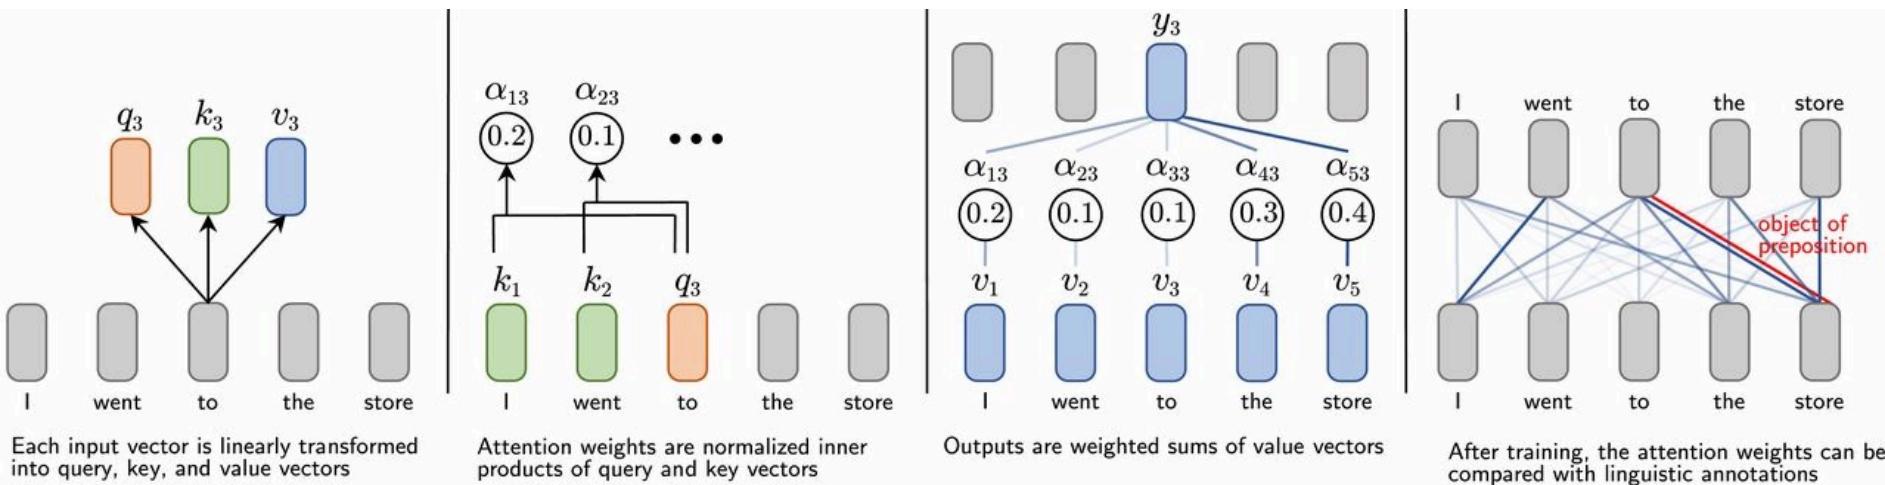
Figure 1: The Transformer - model architecture.



BERT: Bidirectional Encoder Representations from Transformers  
Pre-trained transformer encoder for sentence embedding

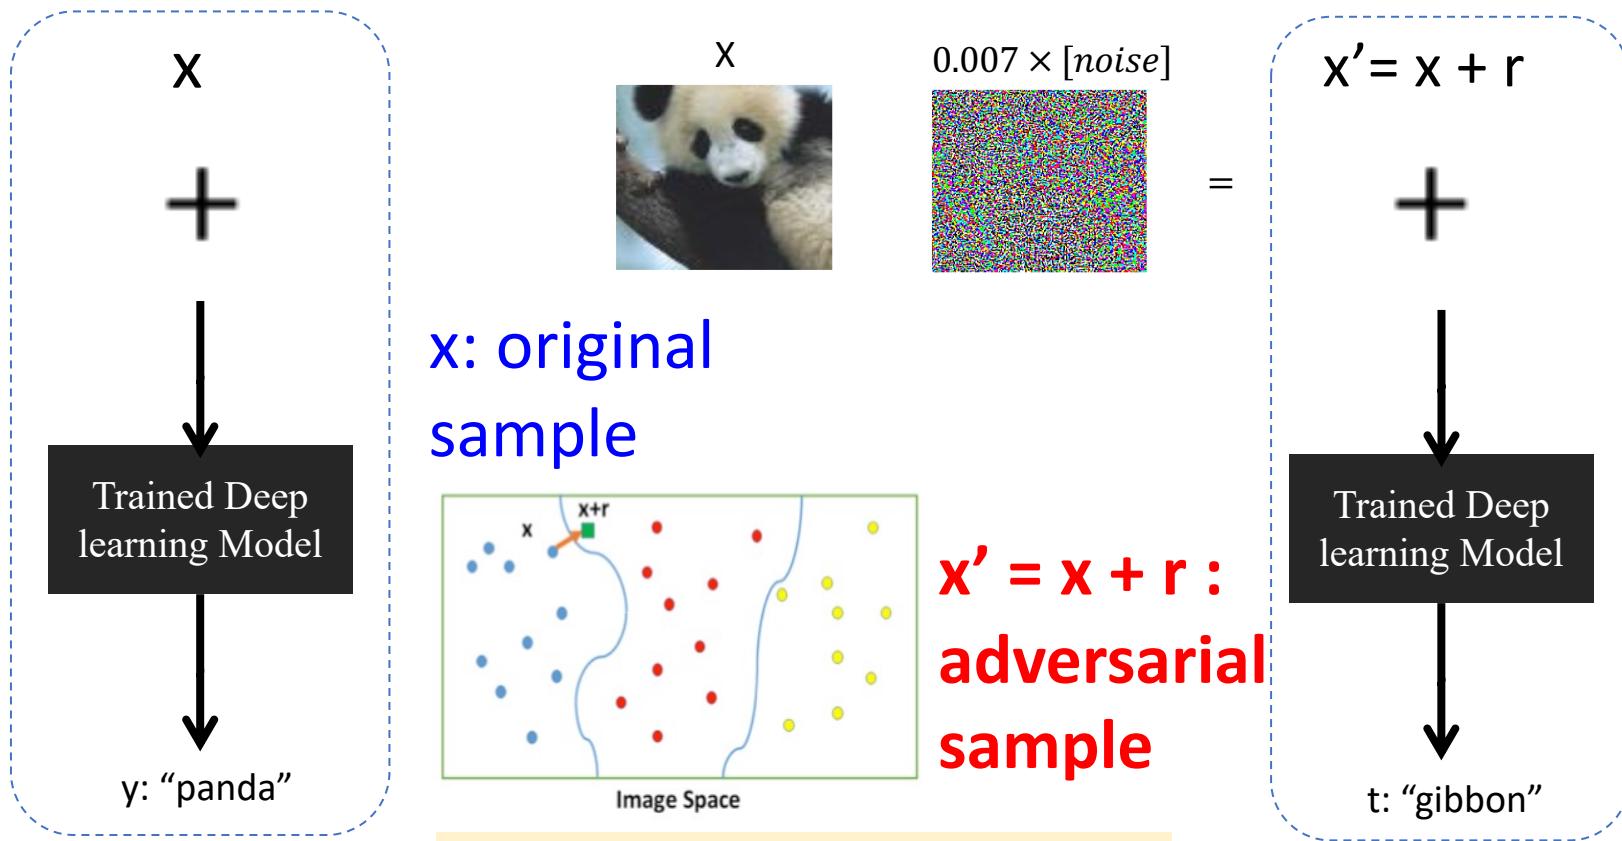


## Notable pre-trained NLP models

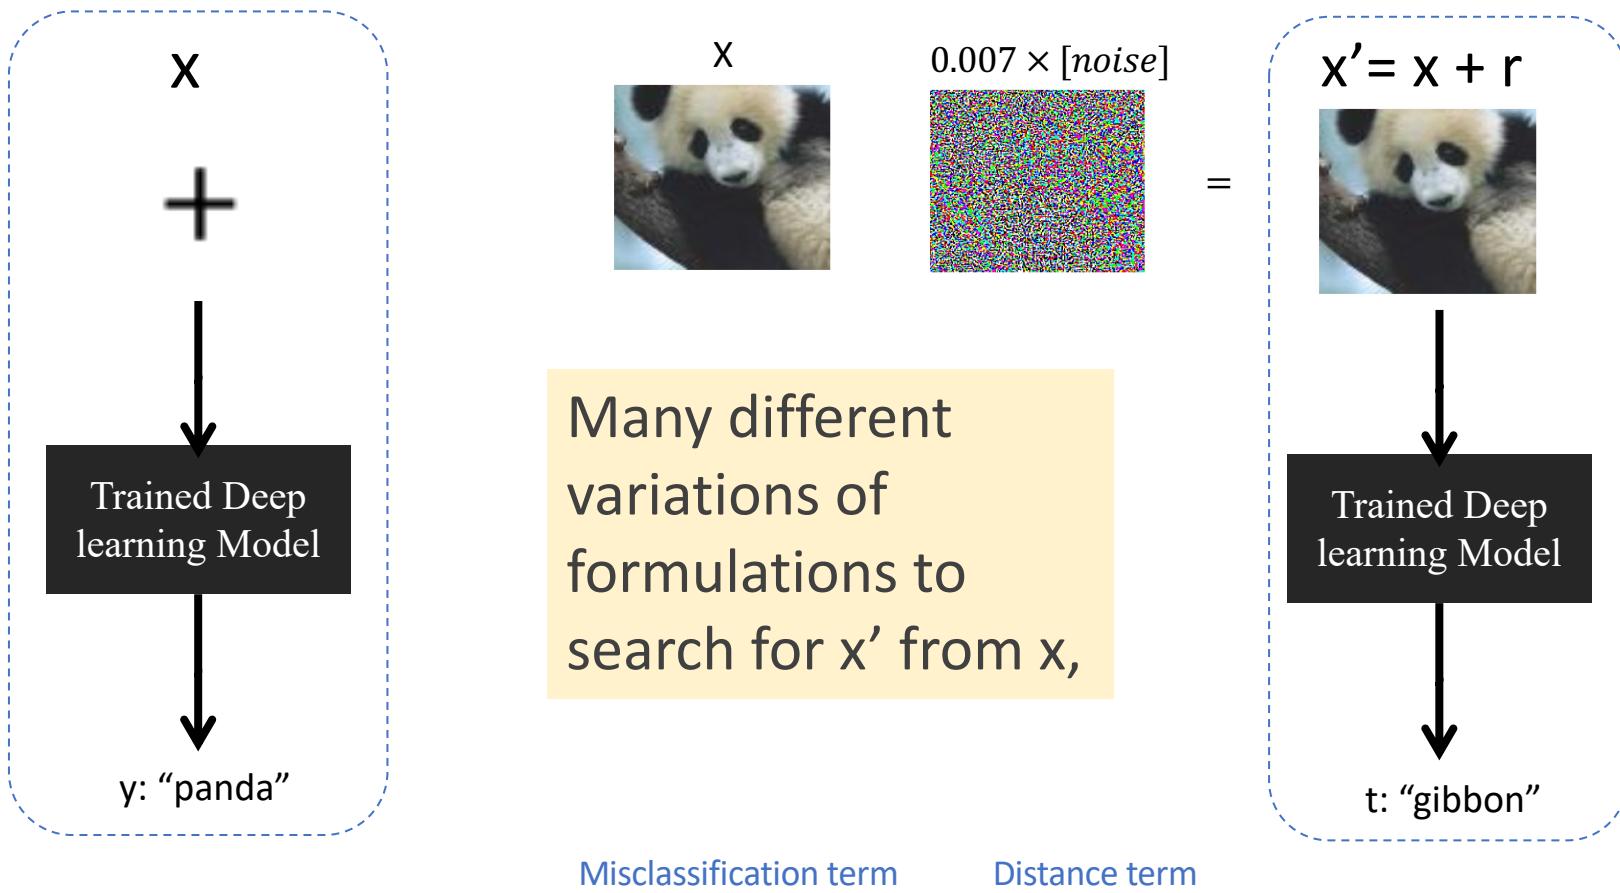


# Background: Adversarial Examples

# Background: Adversarial Examples



# Background: Adversarial Examples



$$\text{minimize } \|f(x') - t\| + \lambda * \Delta(x, x')$$

Misclassification term

Distance term

$$\text{minimize} \|f(x') - t\| + \lambda * \Delta(x, x')$$

# Deep Learning Classifiers are Easily Fooled

## Melanoma Diagnosis with Computer Vision



Healthcare

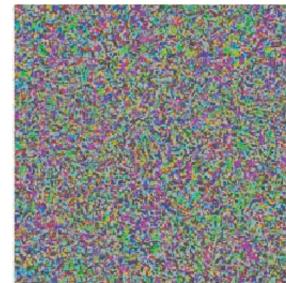
Original Image



Benign

Perturbation

+ 0.04 ×



Adversarial Example

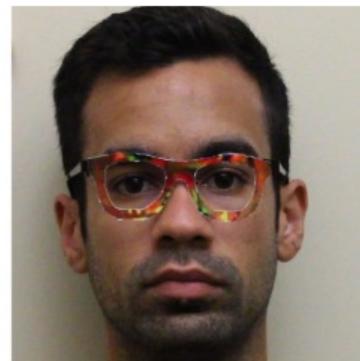
=



Malignant

Samuel G Finlayson et al. "Adversarial attacks on medical machine learning", *Science*, 2019.

# Classifiers Under Attack: Adversary Adapts



**Accessorize to a Crime: Real and Stealthy Attacks on State-of-the-Art Face Recognition**

Mahmood Sharif  
Carnegie Mellon University  
Pittsburgh, PA, USA  
mahmoods@cmu.edu

Sruti Bhagavatula  
Carnegie Mellon University  
Pittsburgh, PA, USA  
srutib@cmu.edu

Michael K. Reiter  
University of North Carolina  
Chapel Hill, NC, USA  
reiter@cs.unc.edu

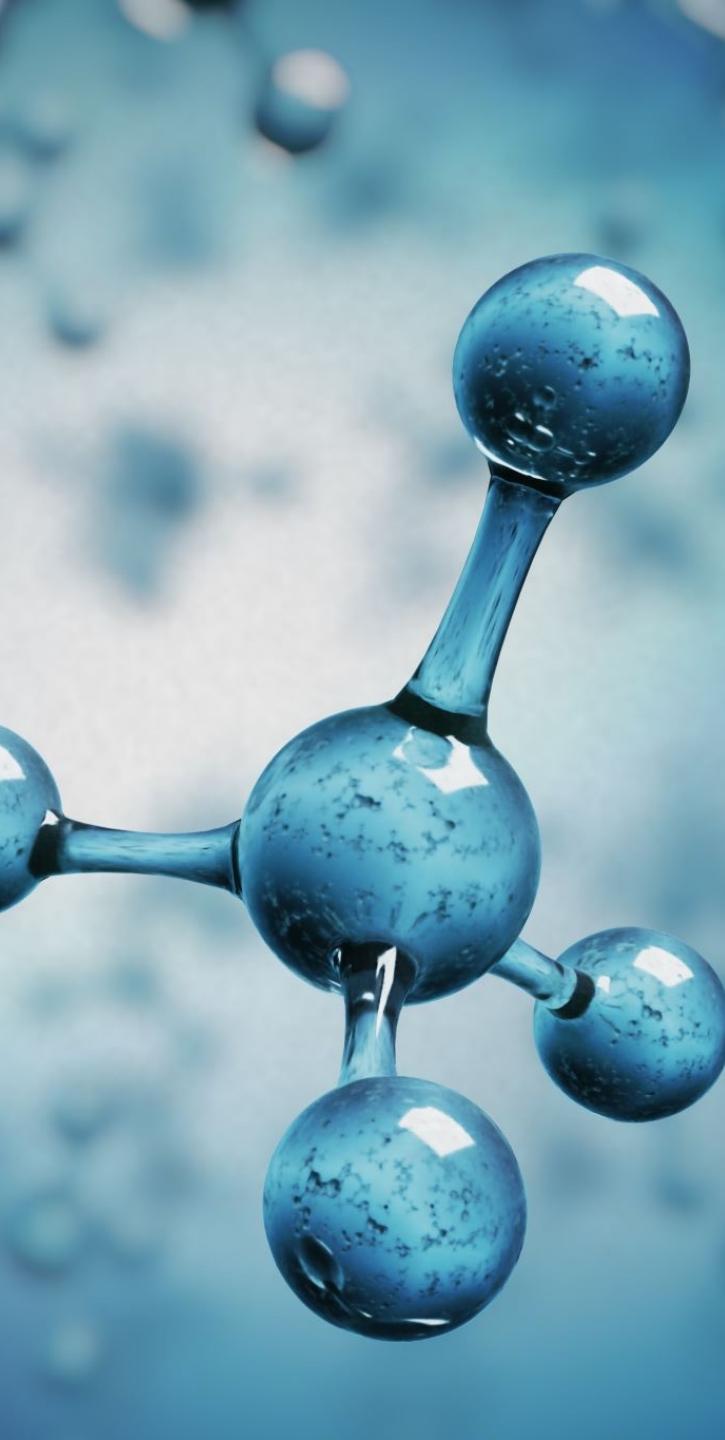
Luoj Bauer  
Carnegie Mellon University  
Pittsburgh, PA, USA  
lbauer@cmu.edu

ACM CCS 2016

Actual images

Recognized faces

Mahmood Sharif et al. "Accessorize to a Crime: Real and Stealthy Attacks on State-of-the-Art Face Recognition", In CCS, 2016.



# Terminology

- Changes to inputs that fool the model are known as **adversarial examples** or **adversarial perturbations**
- A program that repeatedly generates adversarial examples for some model is known as an **adversarial attack**
- A model's resistance to adversarial examples is known as **robustness**

## Toxicity Identification



## Sentiment Classification

"I love this movie.  
I've seen it many times  
and it's still awesome."



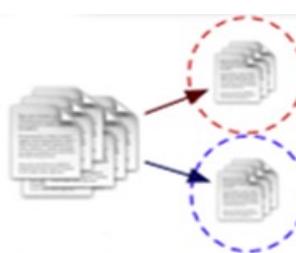
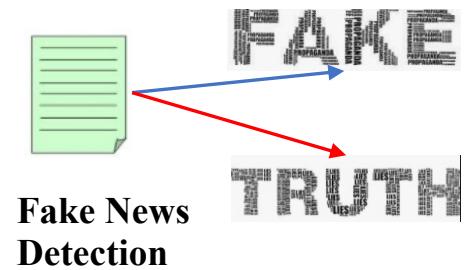
"This movie is bad.  
I don't like it at all.  
It's terrible."

## Authorship Detection



Rowling?

NLP Computer System needs Trustworthiness and Robustness



Electronic Medical Records

Misclassification term

Distance term

$$\text{minimize } \|f(x') - t\| + \lambda * \Delta(x, x')$$

## What are adversarial examples in NLP?

- **Idea 1:** examples that are *almost* visually indistinguishable to humans (**mispellings**)

**Input,  $x$ :**

“True Grit” was the best **movie**  
I’ve seen **since** I was a small boy.

Prediction: **Positive ✓**

**Perturbation,  $x_{\text{adv}}$ :**

“True Grit” was the best **moive**  
I’ve seen **snice** I was a small boy.

Prediction: **Negative X**

Useful, but easy to defend against:

- Pass inputs into a **spell-checker** before feeding them into the model
- Or, train an RNN to correct inputs

25

Misclassification term

Distance term

$$\text{minimize } \|f(x') - t\| + \lambda * \Delta(x, x')$$

## What are adversarial examples in NLP?

- **Idea 2:** examples that are indistinguishable *in meaning* to the original input (**semantics-preserving changes**)

**Input,  $x$ :**

“True Grit” was the best movie  
I’ve seen since I was a **small boy**.

Prediction: **Positive ✓**

**Perturbation,  $x_{\text{adv}}$ :**

“True Grit” was the best movie  
I’ve seen since I was a **wee lad**.

Prediction: **Negative X**



AE NLP  
literature  
is messy  
(chaotic)

1. Many generated examples are bad

2. No standard library

3. No clear benchmarking insights

4. No clear benefits

Our Solution:  
TextAttack to Rescue

1. Many generate  
examples are bad

AE NLP  
literature  
is messy  
(chaotic)

### Input, $x$ :

“True Grit” was the best movie I’ve seen since I was a **small boy**.

Prediction: **Positive ✓**

### Perturbation, $x_{\text{adv}}$ :

“True Grit” was the best movie I’ve seen since I was a **wee lad**.

Prediction: **Negative X**

## Bad examples of adversarial perturbations in NLP

### Perturbation, $x_{\text{adv}}$ :

different **semantics** than original input



“True Grit” was the **worst** movie I’ve seen since I was a small boy.

violates **grammar** (unlike the original input)



“True Grit” was the best movie I’ve seen since I **were boy small**.

this is just **suspicious** – nobody talks like that!



“True Grit” was the best movie I’ve seen since I was a **minuscule youngster**.

# Constraints to ensure our transformation only produces “valid” examples?

- **Idea 1:** what is the cosine similarity between the sentence embeddings of  $\mathbf{x}$  and  $\mathbf{x}_{\text{adv}}$ ?
  - (we can obtain sentence embeddings from the Universal Sentence Encoder, for example)
- **Idea 2:** Use a grammar checker to sure that we didn't introduce any grammatical errors in  $\mathbf{x}_{\text{adv}}$ .

Let  $T(x)$  be transformation and  $C_i(x)$  be a constraint,

$$C_1(T(x)) \wedge C_2(T(x)) \wedge \dots \wedge C_m(T(x))$$

31

# Standardize Constraints Enables Better/Truthful Comparison

Constraints Search Method	TFADJUSTED		TEXTFOOLER	
	TEXTFOOLER	GENETICATTACK	TEXTFOOLER	GENETICATTACK
Semantic Preservation	4.06	4.11	-	-
Grammatical Error %	0	0	-	-
Non-suspicion Score	58.8	56.9	-	-
Attack Success %	10.6	<b>12.0</b>	91.1	<b>95.0</b>
Perturbed Word %	11.1	11.0	18.9	17.2
Num Queries	<b>27.1</b>	4431.6	<b>77.0</b>	3225.7

Table 7: Comparison of the search methods from GENETICATTACK and TEXTFOOLER with two sets of constraints (TEXTFOOLER and TFADJUSTED). Attacks were run on 1000 samples against BERT fine-tuned on the MR dataset. GENETICATTACK’s genetic algorithm is more successful than TEXTFOOLER’s greedy strategy, albeit much less efficient.

Our Analysis paper: Reevaluating Adversarial Examples in Natural Language

- 2020 [EMNLP Findings](#)

AE NLP  
literature  
is messy  
(chaotic)

2. No standard library

# Problems with Current NLP Attack Ecosystem

Many attacks, but Each implemented and benchmarked in separate codebases (if released at all)

- Hard to trust literature comparisons because implementation differences can affect results
- hard to benchmark

Challenging to develop new attacks re-using existing components

- Lots of overlap between attacks (e.g. synonym substitution techniques), but little standardization or re-usability

Difficult to utilize attacks and attack components for improving models

- Attack implementations are almost never model-agnostic
- Adversarial training code is usually unreleased or non-existent
- Data augmentation not nearly as commonplace as in images

# Generating NLP adversarial examples

## Four Components Framework:

1. Goal Function: defines end-goal for adversarial attack
2. Constraints: linguistic requirements for valid adversarial examples
3. Transformation: mechanism for generating potential adversarial examples
4. Search Algorithm: method for finding sequence of transformations that produce valid adversarial examples defined by goal function and constraints

Goal Function term	Constraints' term
$\text{minimize } \ f(x') - t\  + \lambda * \Delta(x, x')$	

Tool Paper: **TextAttack: A Framework for Adversarial Attacks, Data Augmentation, and Adversarial Training in NLP**  
•2020 [EMNLP Demo](#)

# Transformation: Word Substitution centered

- **Thesaurus:** Look up the word in a thesaurus
- **Embeddings:** Search for nearest-neighbors in the embedding space
- **Hybrid:** Search for nearest neighbors in the *counter-fitted* embedding space (*Mrkšić et al, 2016*)

all of these are TextAttack **transformations**  
(textattack.transformations)

# How can we use transformations and constraints to attack a NLP model?

- We need two more things:
  - 1. A way to **search** the space of transformations for a valid, successful adversarial example.
  - 2. A way to know whether an example successfully fools the model.

TextAttack **goal functions**  
(textattack.goal\_functions)

TextAttack **search methods**  
(textattack.search\_methods)

Goal Function term

Constraints' term

$$\text{minimize } \|f(x') - t\| + \lambda * \Delta(x, x')$$

# The TextAttack Framework

NLP attacks can be constructed from four components:

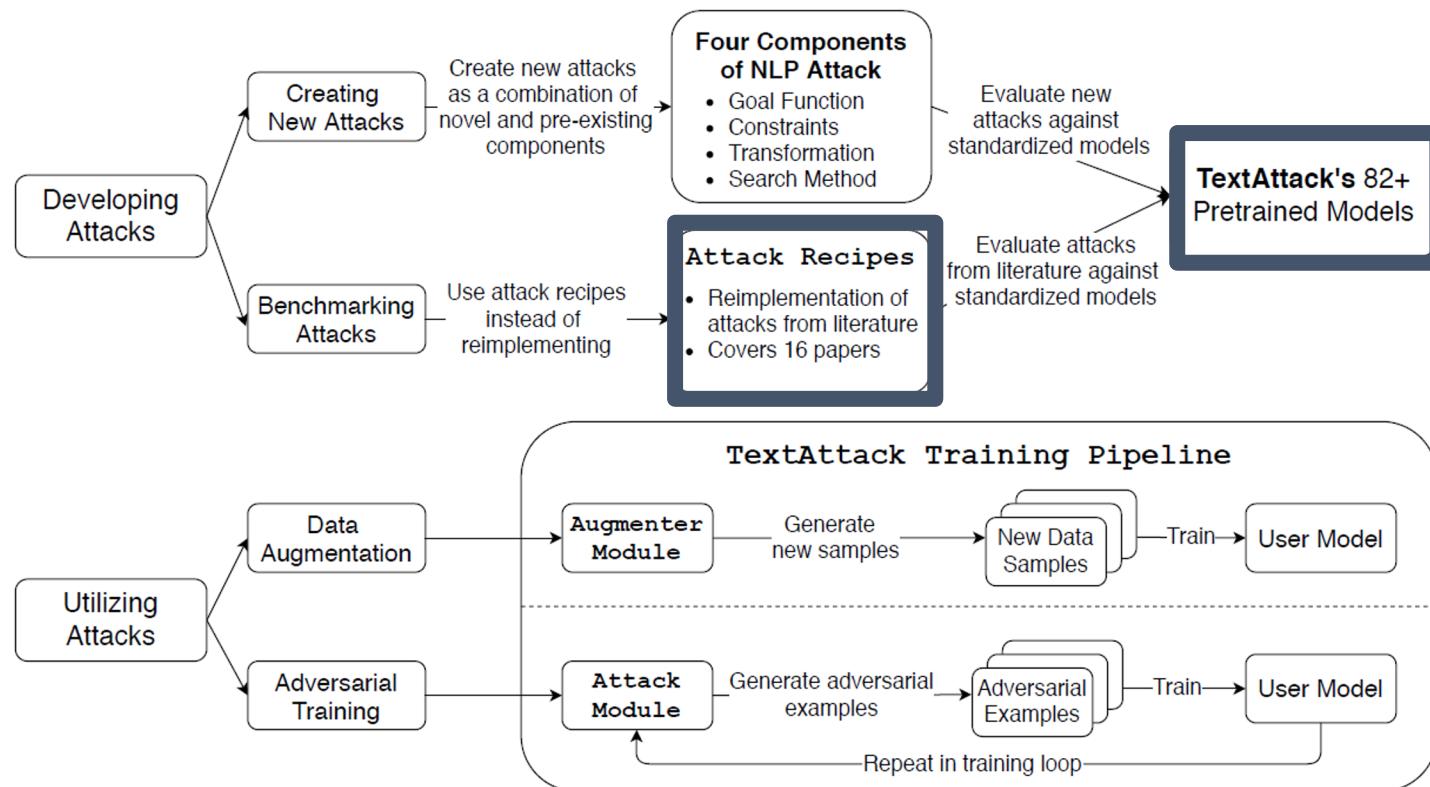
- 1. transformation** ([textattack.transformations.Transformation](#))
- 2. constraint(s)** ([list\(textattack.constraints.Constraint\)](#))
- 3. goal function** ([textattack.goal\\_functions.GoalFunction](#))
- 4. search method** ([textattack.search\\_methods.SearchMethod](#))

Goal Function term

Constraints' term

$$\text{minimize } \|f(x') - t\| + \lambda * \Delta(x, x')$$

# TextAttack's Features



# Is BERT Really Robust? A Strong Baseline for Natural Language Attack on Text Classification and Entailment

Di Jin,<sup>1\*</sup> Zhijing Jin,<sup>2\*</sup> Joey Tianyi Zhou,<sup>3</sup> Peter Szolovits<sup>1</sup>

<sup>1</sup>Computer Science & Artificial Intelligence Laboratory, MIT

<sup>2</sup>University of Hong Kong

<sup>3</sup>A\*STAR, Singapore

jindi15@mit.edu, zhijing.jin@connect.hku.hk, zhouty@ihpc.a-star.edu.sg, psz@mit.edu

## Abstract

Machine learning algorithms are often vulnerable to adversarial examples that have imperceptible alterations from the original counterparts but can fool the state-of-the-art models. It is helpful to evaluate or even improve the robustness of these models by exposing the maliciously crafted adversarial examples. In this paper, we present **TEXTFOOLER**, a simple but strong baseline to generate adversarial text. By applying it to two fundamental natural language tasks, text classification and textual entailment, we successfully attacked three target models, including the powerful pre-trained BERT, and the widely used convolutional and recurrent neural networks. We demonstrate three advantages of this framework: (1) effective—it outperforms previous attacks by success rate and perturbation rate, (2) utility-preserving—it preserves semantic content, grammaticality, and correct types classified by humans, and (3) efficient—it generates adversarial text with computational complexity linear to the text length.<sup>1</sup>

## Classification Task: Is this a *positive* or *negative* review?

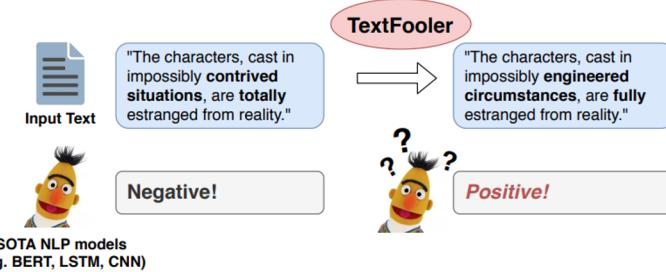


Figure 1: Our model TextFooler slightly change the input text but completely altered the prediction result.

al. 2013; Carlini and Wagner 2018), it is still challenging to deal with text data due to its discrete nature. Formally, besides the ability to fool the target models, outputs of a natural

---

**Algorithm 1** Adversarial Attack by TEXTFOOLER

---

**Input:** Sentence example  $X = \{w_1, w_2, \dots, w_n\}$ , the corresponding ground truth label  $Y$ , target model  $F$ , sentence similarity function  $\text{Sim}(\cdot)$ , sentence similarity threshold  $\epsilon$ , word embeddings  $\text{Emb}$  over the vocabulary  $\text{Vocab}$ .

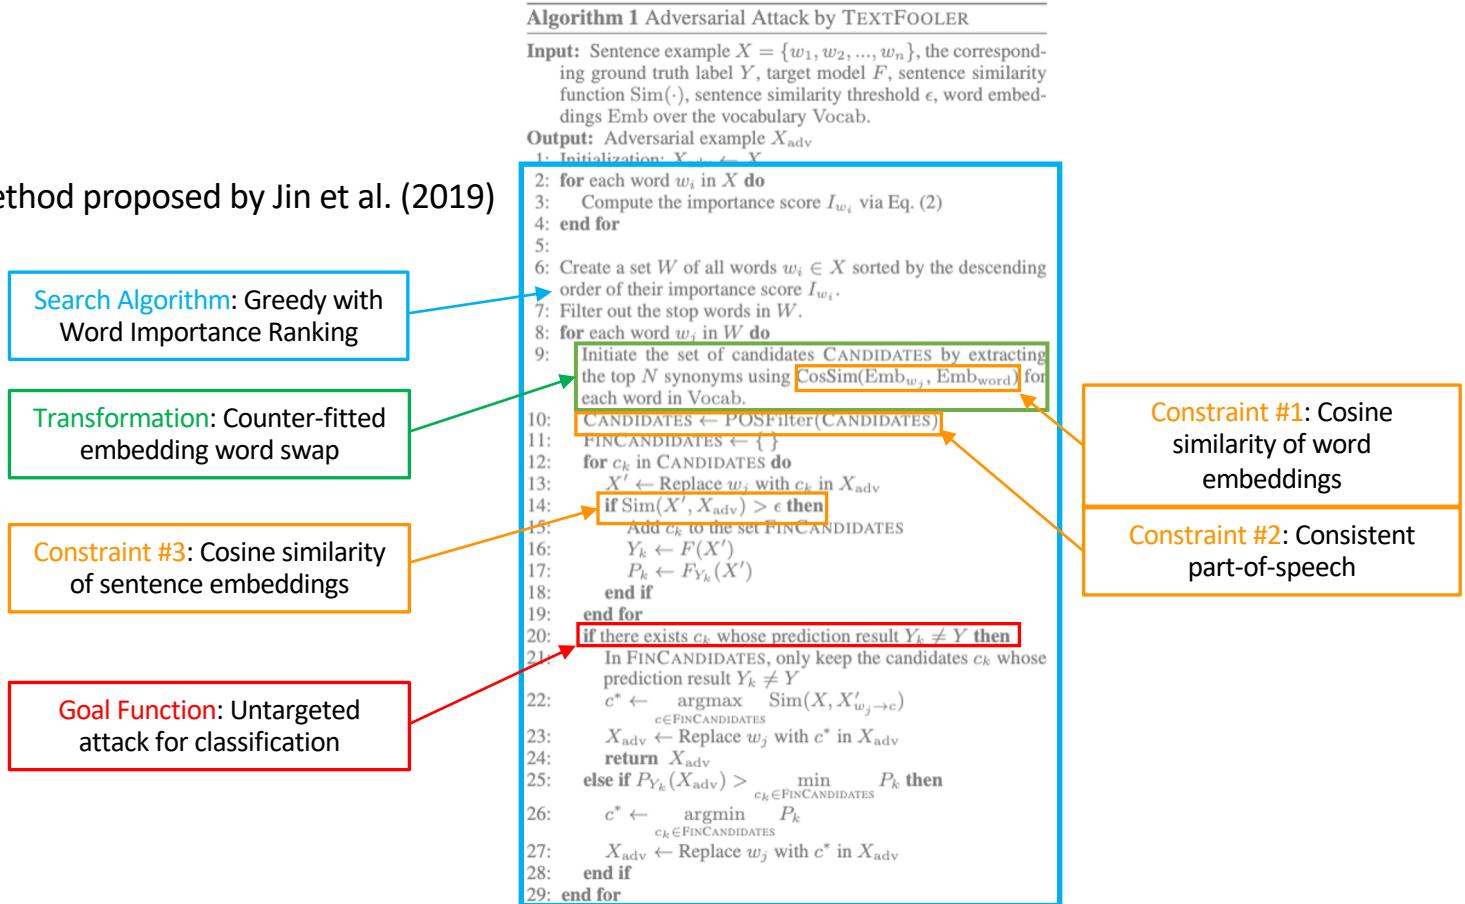
**Output:** Adversarial example  $X_{\text{adv}}$

- 1: Initialization:  $X_{\text{adv}} \leftarrow X$
- 2: **for** each word  $w_i$  in  $X$  **do**
- 3:   Compute the importance score  $I_{w_i}$  via Eq. (2)
- 4: **end for**
- 5:
- 6: Create a set  $W$  of all words  $w_i \in X$  sorted by the descending order of their importance score  $I_{w_i}$ .
- 7: Filter out the stop words in  $W$ .
- 8: **for** each word  $w_j$  in  $W$  **do**
- 9:   Initiate the set of candidates  $\text{CANDIDATES}$  by extracting the top  $N$  synonyms using  $\text{CosSim}(\text{Emb}_{w_j}, \text{Emb}_{\text{word}})$  for each word in  $\text{Vocab}$ .
- 10:    $\text{CANDIDATES} \leftarrow \text{POSFilter}(\text{CANDIDATES})$
- 11:    $\text{FINCANDIDATES} \leftarrow \{\}$
- 12:   **for**  $c_k$  in  $\text{CANDIDATES}$  **do**
- 13:      $X' \leftarrow \text{Replace } w_j \text{ with } c_k \text{ in } X_{\text{adv}}$
- 14:     **if**  $\text{Sim}(X', X_{\text{adv}}) > \epsilon$  **then**
- 15:       Add  $c_k$  to the set  $\text{FINCANDIDATES}$
- 16:        $Y_k \leftarrow F(X')$
- 17:        $P_k \leftarrow F_{Y_k}(X')$
- 18:     **end if**
- 19:   **end for**
- 20:   **if** there exists  $c_k$  whose prediction result  $Y_k \neq Y$  **then**
- 21:     In  $\text{FINCANDIDATES}$ , only keep the candidates  $c_k$  whose prediction result  $Y_k \neq Y$
- 22:      $c^* \leftarrow \underset{c \in \text{FINCANDIDATES}}{\text{argmax}} \text{Sim}(X, X'_{w_j \rightarrow c})$
- 23:      $X_{\text{adv}} \leftarrow \text{Replace } w_j \text{ with } c^* \text{ in } X_{\text{adv}}$
- 24:     **return**  $X_{\text{adv}}$
- 25:   **else if**  $P_{Y_k}(X_{\text{adv}}) > \min_{c_k \in \text{FINCANDIDATES}} P_k$  **then**
- 26:      $c^* \leftarrow \underset{c_k \in \text{FINCANDIDATES}}{\text{argmin}} P_k$
- 27:      $X_{\text{adv}} \leftarrow \text{Replace } w_j \text{ with } c^* \text{ in } X_{\text{adv}}$
- 28:   **end if**
- 29: **end for**
- 30: **return** None

---

# Four Components in Action

TextFooler method proposed by Jin et al. (2019)



# Four Components Standardized 18 Attacks:

	Alzantot et al. (2018)	Jin et al. (2019)
Goal Function	UntargetedClassification	UntargetedClassification
Search Method	GeneticAlgorithmWordSwap	GreedyWordSwapWordImportanceRanking
Transformation	WordSwapEmbedding(embedding='cf')	WordSwapEmbedding(embedding='cf')
Constraints	<ul style="list-style-type: none"><li>WordsPerturbedPercentage(max_perc=20)</li><li>WordEmbeddingDistance(max_mse=0.5)</li><li>GoogleLanguageModel(n_per_index=4)</li></ul>	<ul style="list-style-type: none"><li>WordEmbeddingDistance(min_cos_sim=0.5)</li><li>PartOfSpeech(verb_noun_swap=True)</li><li>UniversalSentenceEncoder(metric='angular', thresh=0.904458599)</li></ul>

# Pretrained Models



Integration with  
HuggingFace's [Model Hub](#)  
and [nlp](#) library

Can attack any model on the  
model hub on any dataset from  
[nlp](#)



TextAttack has 82 pretrained  
models on its [Model Hub](#)  
[page](#)

Models: BERT, DistilBERT,  
ALBERT, BART, RoBERTa, XLNet  
Trained on all [GLUE](#) tasks

# Installing TextAttack

 Github PyTest passing pypi package 0.2.12

**pip install textattack**

<https://github.com/QData/TextAttack>

[QData / TextAttack](#)

Code Issues 35 Pull requests 9 Actions Projects 6 Wiki Security Insights Settings

master 22 branches 9 tags Go to file Add file Code

 qiyanjun Update README.md ... ✓ ae68c81 5 days ago 1,989 commits

File	Description	Time
.github	Update run-pytest.yml	2 months ago
docs	Fix errors in Example_5_Explain_BERT	6 days ago
examples	isort format of attack_camerbert	6 months ago
tests	Revert "add --split to specify train/test/dev dataset"	12 days ago
textattack	Revert "add --split to specify train/test/dev dataset"	12 days ago
.gitignore	delete vscode setting	7 months ago
.readthedocs.yml	fix readthedocs module load	5 months ago
CONTRIBUTING.md	Clarify CONTRIBUTING.md	8 months ago
LICENSE	Initial commit	2 years ago
Makefile	autobuild cli changed	13 days ago
README.md	Update README.md	5 days ago
README_ZH.md	correct the EMNLP BlackBoxNLP mentions.	4 months ago
pytest.ini	update travis for jenkins	10 months ago
requirements.txt	locally test all passed...	9 days ago
setup.cfg	merge in master and fix syntax errors	10 months ago
setup.py	Update setup.py	9 days ago

About

TextAttack 🤖 is a Python framework for adversarial attacks, data augmentation, and model training in NLP

🔗 [textattack.readthedocs.io/en/latest/](https://textattack.readthedocs.io/en/latest/)

nlp security machine-learning  
natural-language-processing  
data-augmentation  
adversarial-machine-learning  
adversarial-examples  
adversarial-attacks

📄 Readme

⚖️ MIT License

Releases 9

↳ v0.2.15: CLARE Attack, Cu... Latest on Dec 26, 2020 + 8 releases

Packages

No packages published

AE NLP  
literature  
is messy  
(chaotic)

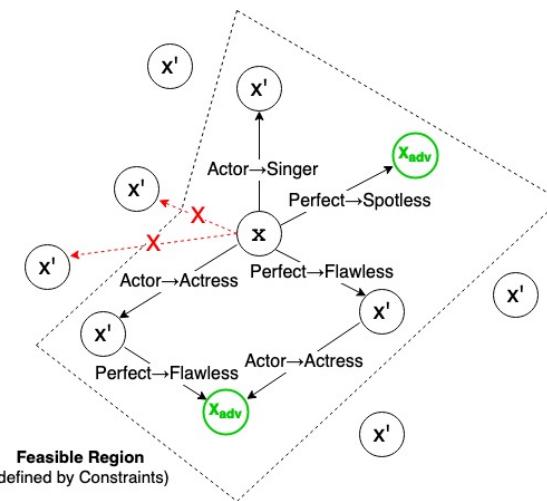
3. No clear benchmarking  
insights

# Search Algorithm

Search Algorithm: A way to **search** the space of transformations for a valid, successful adversarial example.

## Why a search algorithm?

- We need to find set of transformations that successfully produce  $x_{adv}$
- Combinatorial search problem with heuristic  $score(x)$  provided by goal function



Our Analysis paper: Searching for a Search Method: Benchmarking  
Search Algorithms for Generating NLP Adversarial Examples  
•2020 [EMNLP BlackBoxNLP](#)

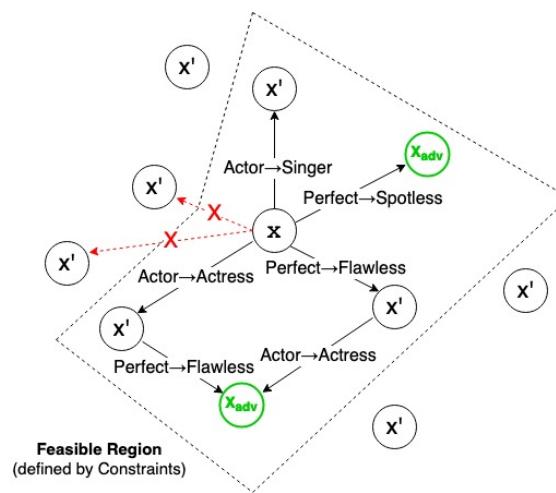
# Search Space

Search Algorithm: A way to **search** the space of transformations for a valid, successful adversarial example.

Search space defined by transformation and constraints

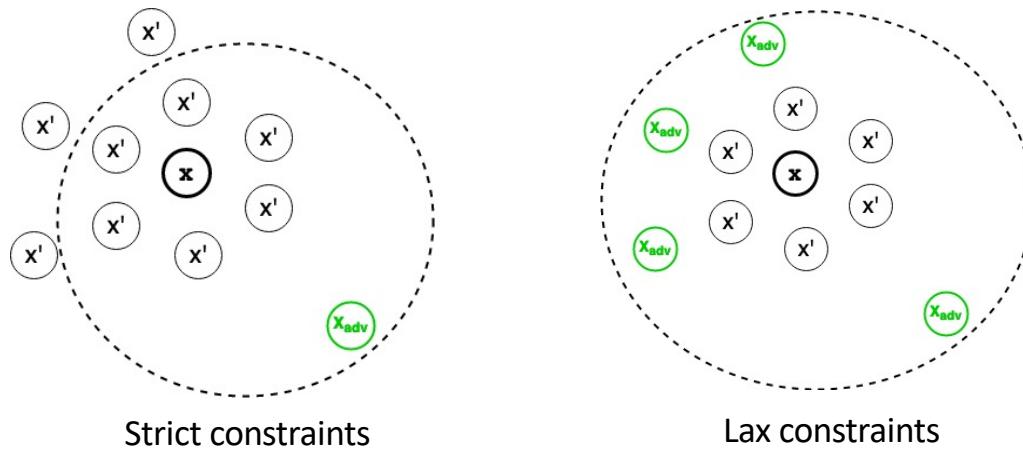
Let  $T(x)$  be our transformation and  $C_i(x)$  be a constraint,

$$S(x) = \{T(x) | C_1(T(x)) \wedge C_2(T(x)) \wedge \dots \wedge C_m(T(x))\}$$



# Search Space

How search space is defined can affect performance of the search algorithm



# Search Algorithms from Literature

A lot of works have proposed novel search algorithms.

Proposed search algorithms:

- **Greedy:** (Kuleshov et al. 2018)
- **Beam Search:** (Ebrahimi et al., 2017)
- **Greedy with Word Importance Ranking:** (Gao et al., 2018), (Jin et al., 2019), (Ren et al., 2019)
- **Genetic Algorithm:** (Alzantot et al., 2018),
- **Particle Swarm Optimization:** (Zang et al., 2020)
- **MCMC Sampling:** (Zhang et al., 2019)

# Problems in Current Literature

Inconsistent search  
space for comparisons

Lack of comprehensive  
*performance*  
benchmark for search  
algorithm

Lack of comprehensive  
*speed* benchmark for  
search algorithm

# Performance across different search methods

Model	Dataset	Search Method	Lax Constraint: $sim = 0.5$			Strict Constraint: $sim = 0.9$		
			A.S. %	P.W. %	Queries	A.S. %	P.W. %	Queries
BERT	Yelp	Word Importance Ranking (UNK)	99.8	8.59	393	25.7	10.69	219
		Word Importance Ranking (DEL)	99.7	9.16	423	26.1	10.73	220
		Word Importance Ranking (RAND)	99.7	16.43	610	23.5	12.58	94
		Greedy (b=1)	99.8	5.02	9,813	30.3	7.59	1,984
		Beam Search (b=4)	100.0	4.92	30,417	31.1	7.59	7,297
		Beam Search (b=8)	100.0	4.89	57,984	31.3	7.59	14,329
		Genetic Algorithm	99.6	9.83	7,173	21.5	9.52	12,655
BERT	MR	Word Importance Ranking (UNK)	99.2	15.58	116	30.2	14.5	34
		Word Importance Ranking (DEL)	98.8	15.00	113	30.8	15.04	34
		Word Importance Ranking (RAND)	99.1	21.00	132	29.1	16.00	16
		Greedy (b=1)	99.3	11.86	639	31.1	11.51	49
		Beam Search (b=4)	99.7	11.68	1,411	32.1	11.64	141
		Beam Search (b=8)	99.7	11.61	2,432	32.3	11.67	261
		Genetic Algorithm	99.4	14.93	1,611	31.4	13.47	2,870
BERT	SNLI	Word Importance Ranking (UNK)	100.0	7.05	66	35.5	10.45	34
		Word Importance Ranking (DEL)	100.0	7.49	68	35.4	10.52	34
		Word Importance Ranking (RAND)	99.9	13.60	89	33.1	12.22	13
		Greedy (b=1)	100.0	6.08	473	38.3	7.94	37
		Beam Search (b=4)	100.0	6.02	662	39.7	7.97	95
		Beam Search (b=8)	100.0	6.02	918	40.1	8.09	173
		Genetic Algorithm	100.0	7.05	996	39.1	9.44	2,332

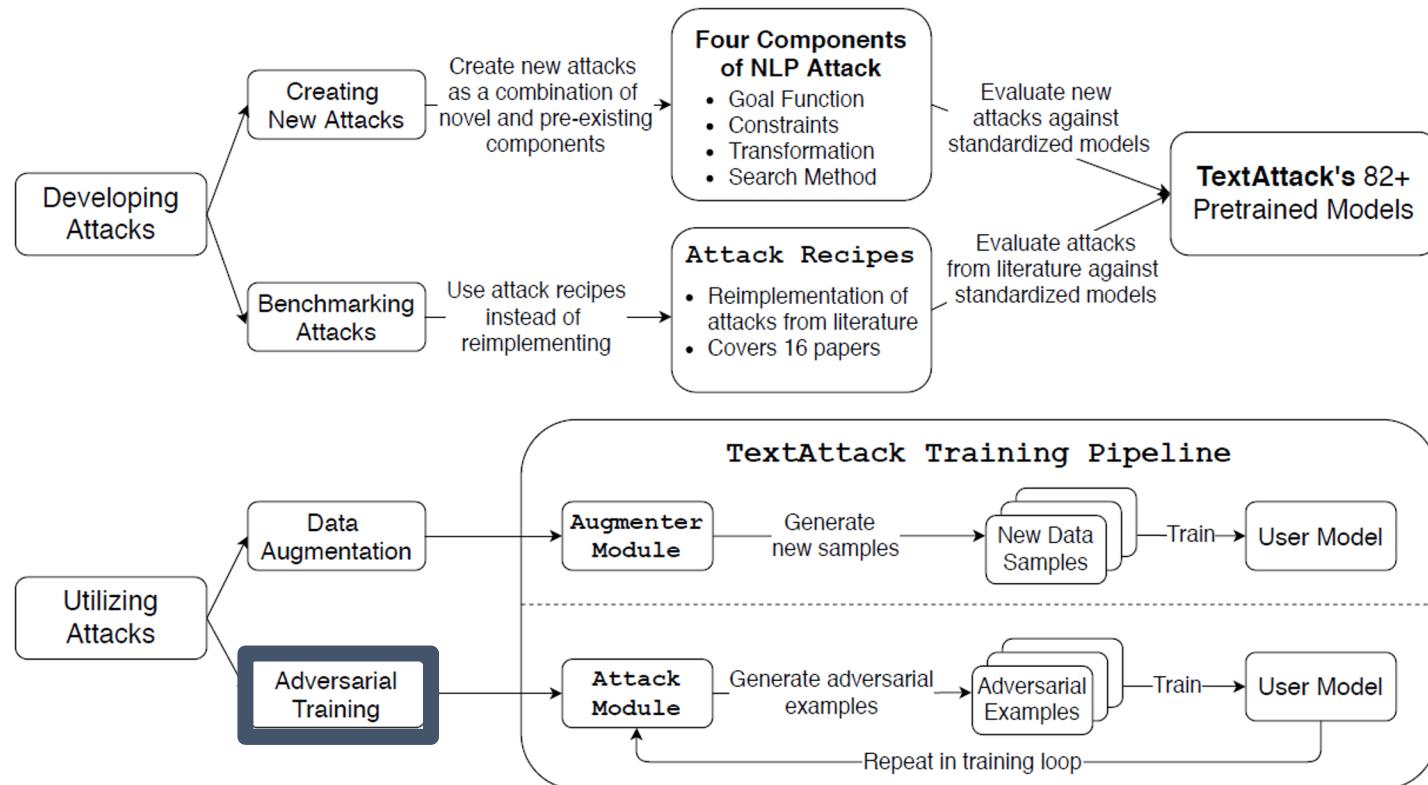
# Benchmarking Insights

- Optimal method for absolute performance is beam search with beam width of 8.
- When within a small query budget, greedy with word importance ranking is most effective
  - For two constraint settings across three datasets, the relative differences between the attack success rates of greedy with word importance ranking and the success rates of beam search are less than 20%.
- Search algorithms matter less than transformations and constraints.
  - Although changing the search methods did not change attack success rate by more than 20%, changing the constraints changed attack success rate by over 60%.

AE NLP  
literature  
is messy  
(chaotic)

4. No clear benefits

# Adversarial Training



# Adversarial Training for Robustness

Goodfellow et al. (2015): 
$$\arg \min_{\theta} \left[ \mathbb{E}_{(x,y) \in \hat{p}_{\text{data}}} \left( \max_{\delta \in S} L(\theta, x + \delta, y) \right) + \mathbb{E}_{(x,y) \in \hat{p}_{\text{data}}} \left( L(\theta, x, y) \right) \right] \quad (2)$$

Madry et al. (2017):  $\arg \min_{\theta} \mathbb{E}_{(x,y) \in \hat{p}_{\text{data}}} \left( \max_{\delta \in S} L(\theta, x + \delta, y) \right)$  (1)

$$J(\mathbb{M}, \boldsymbol{\theta}) + \lambda \frac{1}{m} \sum_{i=1}^m L\left(f(\mathbf{x}^{(i)}; \boldsymbol{\theta}), f(\tilde{\mathbf{x}}^{(i)}; \boldsymbol{\theta})\right).$$

Kannan et al. (2018):

The diagram consists of two parallel vertical arrows pointing upwards. The left arrow points from the text "Adversarial loss" to the label "i=1". The right arrow points from the text "Adversarial Logit Pairing" to the same label "i=1".

# IMDB-BERT Results

	IMDB Test Acc	Yelp Test Acc	Counterfactual Acc
No adv. training	93.97	92.86	92.84
SSMBA	93.94	92.52	92.48
Backtranslation	93.97	92.62	92.58
Textfooler-Mod	<b>94.49</b>	<b>93.29</b>	<b>93.23</b>
BAE-Mod	93.05	91.61	91.35

# TextAttack Rescues Messy AE NLP literature

1. Many generated examples are bad
2. No standard library
3. No clear benchmarking insights
4. No clear benefits

# What can I do with TextAttack?

- build an NLP attacks from a library of components
- run those attacks on models & datasets (yours or ours)
- visualize attack results using the command line, Visdom, W&B, etc.
- or, use the infrastructure of TextAttack to develop and benchmark your own NLP attacks
- or, use the components from TextAttack for **data augmentation**
- or, use the components of TextAttack for **adversarial training**

# Who is TextAttack for?

- **researchers** who want to implement new NLP attacks or compare them in a standardized framework
- **any machine learning practitioner** who want to understand the limitations of NLP models and use adversarial training to make their models better
- **anyone training an NLP model** who wants to apply data augmentation to increase test-set accuracy by 1-2%

# Acknowledgements

**My Students on this project:**



**UVA Computer Science Dept. Security  
Research Group:** Prof. David Evans

**UVA Computer Science Dept. Software Safety  
Group:** Prof. Matthew B Dwyer

**UVA Computer Science Dept. NLP  
Research Group:** Prof. Yangfeng Ji

**UVA Computer Science Dept. Software  
Engineering Group:** Prof. Mary Lou Soffa



4/18/21

Yanjun Qi / UVA CS



61

<http://trustworthymachinelearning.org>

## 1. To Fool / Evade Learned Models

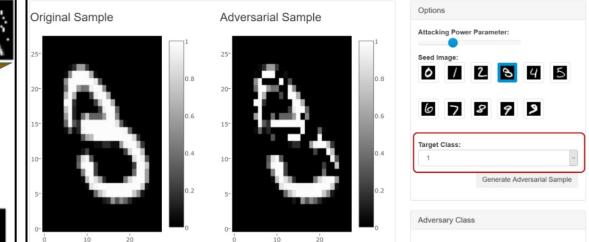
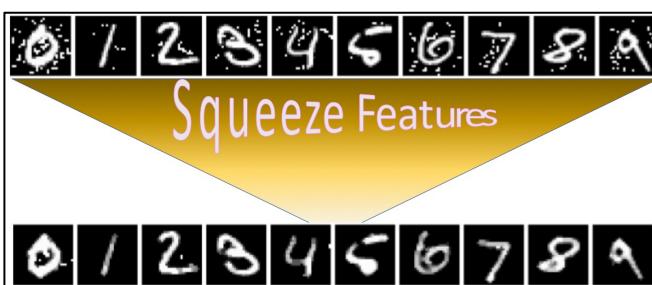
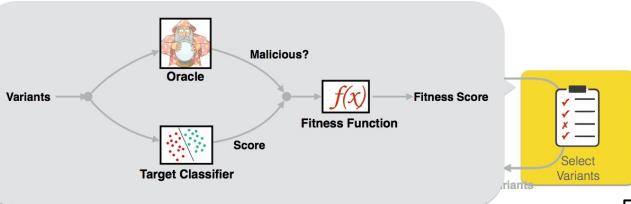
## 2. To Detect fooling/ Evasion

## 3. To Defend Against Evasion

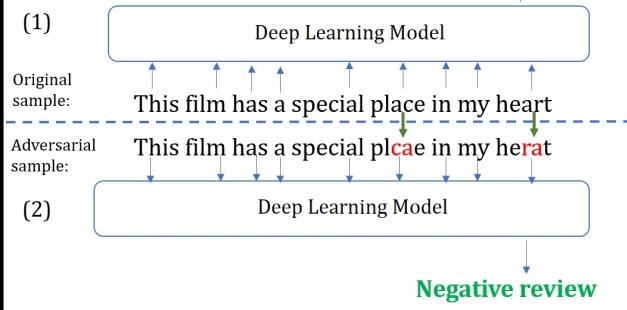
## 4. To Visualize and Benchmarking

## 5. To Understand Theoretically

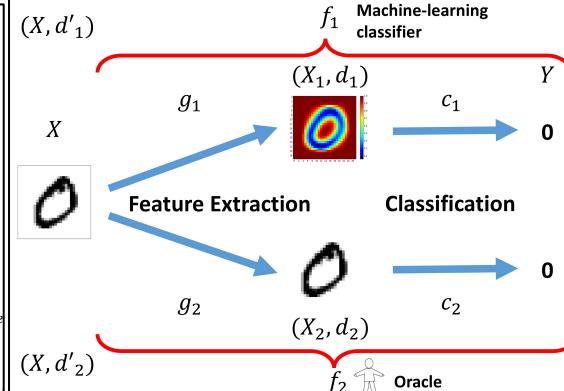
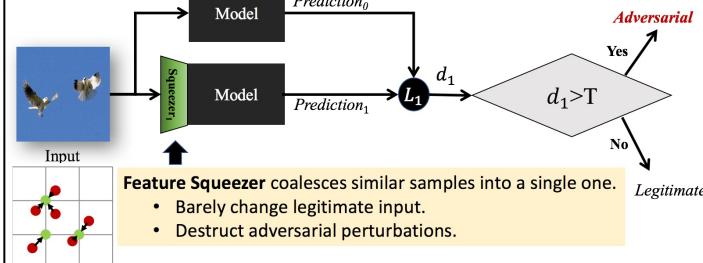
### Automated Evasion Approach Based on Genetic Programming

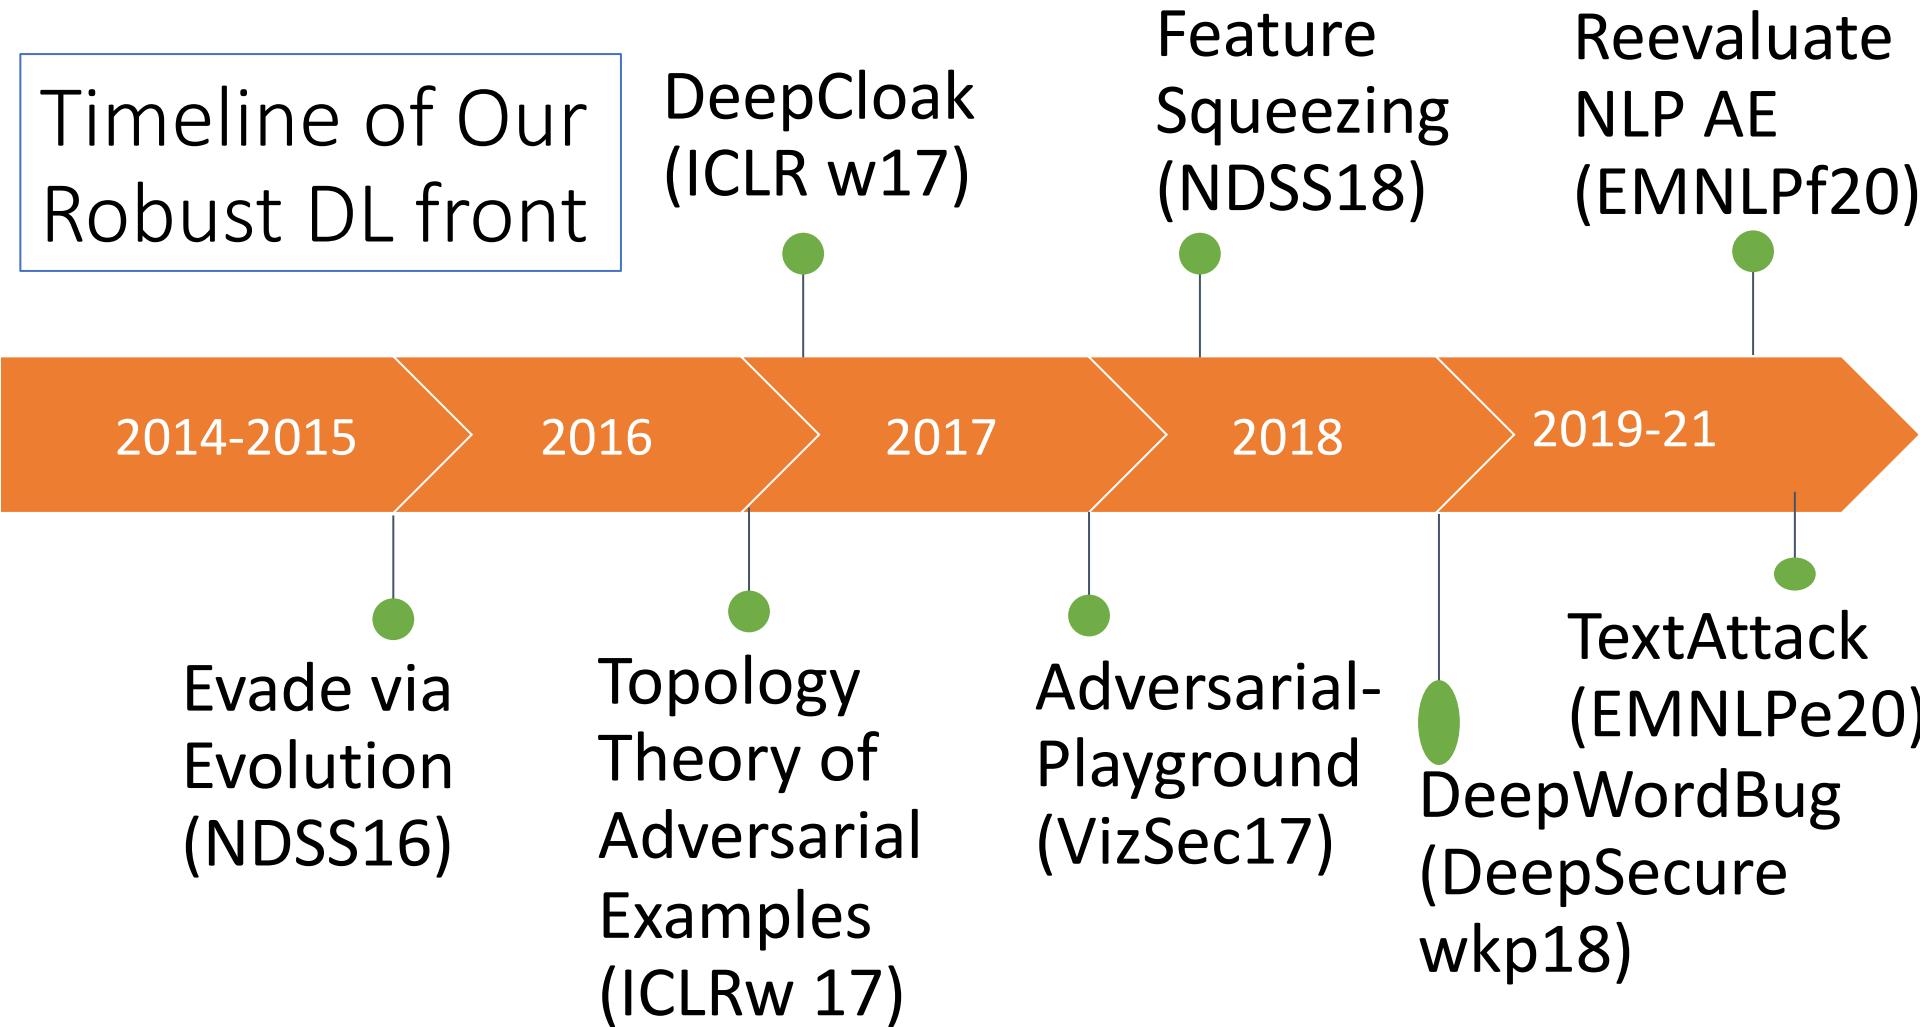


### Positive review



### Detection Framework







Thank you